

Optimization of my_strchr

Original	Optimized
<pre>.global my_strchr my_strchr: src .req r0 c .req r1 i .req r2 t .req r3 ret .req r0 mov i, #0 loop: ldrb t, [src, +i] cmp t, #0 beq loop_end cmp t, c beq found add i, i, #1 b loop found: mov ret, i b end loop_end: mov ret, #-1 b end end: bx lr</pre>	<pre>.global my_strchr_opt my_strchr_opt: src .req r0 c .req r1 i .req r2 t .req r3 ret .req r0 mov i, #0 ldrb t, [src, +i] loop: cmp t, #0 beq loop_end cmp t, c beq found add i, i, #1 ldrb t, [src, +i] b loop found: mov ret, i bx lr loop_end: mov ret, #-1 bx lr</pre>

The existing code has a number of inefficiencies that need to be fixed. First, there is a two cycle stall after the `ldrb` instruction because of the dependency on `t`. Second, there are unnecessary `b end` instructions involved in the end of the routine that can be removed.

To fix the stalls, pre-loading of `t` is used. This means that there will be a 2 cycle stall for the first iteration of the loop, but future iterations will not have a stall. For most loop runs, this will result in a 2 cycle savings. The original loop requires 11 cycles in the normal case, and the optimized code requires only 9 cycles.

To fix the unnecessary `b end` instructions, the `end` label is removed and `bx lr` is repeated (called inlining) in both places where `b end` was previously. This will save 3 cycles as the subroutine finishes, and also removes one instruction from the code.

I decided to apply these changes because they are most straightforward approach to optimizing this code. I also investigated the possibility of applying partial loop unrolling (for example, checking 4 characters per loop iteration) but if you do that then you need to add extra code to the loop to properly calculate `i`, which removes the efficiency gain.