

## **Introduction**

In this report, I will present the optimization of five different string library functions written in ARM assembly. The goal of optimization is to reduce the total number of cycles of runtime, including removing stalls and making use of more efficient approaches to solving the problem. I did not consider code length or registers used in my optimizations, except when those changes can further reduce the number of cycles needed to run.

## strlen

Original	Optimized
<pre>.global my_strlen  my_strlen:     src .req r0     cnt .req r1     t .req r2     ret .req r0      mov cnt, #0 ; Unoptimized loop needs 8 cycles loop:     add cnt, cnt, #1     ldrb t, [src], #1     cmp t, #0     bne loop      sub ret, cnt, #1     bx lr</pre>	<pre>.global my_strlen  my_strlen:     src .req r0     cnt .req r1     t .req r2     ret .req r0      mov cnt, #0 ; t is preloaded     ldrb t, [src], #1 ; Optimized loop needs 6 cycles loop:     add cnt, cnt, #1     cmp t, #0  ; t is preloaded ; This load is also made ; conditional.     ldrbne t, [src], #1     bne loop      sub ret, cnt, #1     bx lr</pre>

The existing code has one main inefficiency that needs to be fixed. There is a two cycle stall after the `ldrb` instruction because of the dependency on `t`.

To fix the stall, pre-loading of `t` is used. This means that there will be a 2 cycle stall for the first iteration of the loop, but future iterations will not have a stall. For most loop runs, this will result in a 2 cycle savings. The original loop requires 8 cycles in the normal case, and the optimized code requires only 6 cycles.

I decided to apply this change because it is the most straightforward approach to optimizing this code. I also investigated the possibility simply moving the `add cnt, cnt, #1` instruction to be between the load of `t` and its usage, but this would only remove 1 stall cycle. Partial loop unrolling could also be used here, potentially for more performance gain. If I wanted to maximize my points, I should probably have done that.

## strcpy

Original	Optimized
<pre>.global my_strcpy  my_strcpy:     dst .req r0     src .req r1     t   .req r2      ; Unoptimized loop needs 8 cycles strcpyloop:     ldrb t, [src], #1     strb t, [dst], #1     cmp t, #0     bne strcpyloop      bx lr</pre>	<pre>.global my_strcpy  my_strcpy:     dst .req r0     src .req r1     t   .req r2      ; t is preloaded     ldrb t, [src], #1      ; Optimized loop needs 6 cycles strcpyloop:     strb t, [dst], #1     cmp t, #0      ;t is preloaded     ;This load is also made     ;conditional.     ldrbne t, [src], #1     bne strcpyloop      bx lr</pre>

The existing code has one main inefficiency that needs to be fixed. There is a stall after the `ldrb` instruction because of the dependency on `t`. How many cycles this stall is, however, is unclear. It is likely two cycles, but because the store instruction doesn't need `t` until the fourth stage in the pipeline, it might only be one cycle. To determine this I would need a cycle accurate simulator, which I don't have access to.

To fix the stall, pre-loading of `t` is used. This means that there will be a stall for the first iteration of the loop, but future iterations will not have a stall. Assuming the stall is 2 cycles, this will result in a 2 cycle savings. The original loop requires 8 cycles in the normal case, and the optimized code requires only 6 cycles.

I decided to apply this change because it is the most straightforward approach to optimizing this code. Partial loop unrolling could also be used here, potentially for more performance gain. If I wanted to maximize my points, I should probably have done that.

## strtolower

Original	Optimized
<pre>.global my_strtolower  my_strtolower:     src .req r0     t   .req r1      ;Number of cycles in the     ;loop is between 12 and 14     ;13 cycles if neither branch (blt     ; and bgt) are executed,     ; 14 cycles if first branch (blt) is     ; not executed, but second branch     ; (bgt) is executed,     ; 12 cycles if first branch (blt) is     ; executed     loop:          ldrb t, [src]         cmp t, #'A'         blt no_match         cmp t, #'Z'         bgt no_match         add t, t, #'a'-'A'     no_match:         strb t, [src], #1         cmp t, #0         bne loop      bx lr</pre>	<pre>.global my_strtolower  my_strtolower:     src .req r0     t   .req r1     tmp .req r2 ; Scratch variable      ;t is preloaded     ldrb t, [src]      ; Optimized loop needs 9 cycles     loop:         ;Branch instructions are         ;removed          ; Compare to test if the         ;letter is uppercase         sub tmp, t, #'A'         cmp tmp, #'Z'-'A'          ;Conditionally make lower         addls t, t, #'a'-'A'          strb t, [src], #1         cmp t, #0          ;t is preloaded         ;This load is also made         ;conditional         ldrbne t, [src]         bne loop      bx lr</pre>

The existing code has a number of inefficiencies that need to be fixed. First, there is a two cycle stall after the `ldrb` instruction because of the dependency on `t`. Second, there are multiple unnecessary branches related to determining if the character is uppercase or not.

To fix the stalls, pre-loading of `t` is used. This means that there will be a 2 cycle stall for the first iteration of the loop, but future iterations will not have a stall. For most loop runs, this will result in a 2 cycle savings. To fix the unnecessary branches, the branch structure is removed and conditional execution is added. The number of cycles saved depends on the character being tested. See the comments in the code for more detail. In the best case, 5 cycles are saved per loop. In the worst case, 3 cycles are saved.

I decided to apply these changes because they are most straightforward approach to optimizing this code. I also investigated the possibility of applying partial loop unrolling, but I think it makes this code needlessly complex.

## strchr

Original	Optimized
<pre>.global my_strchr  my_strchr:     src .req r0     c   .req r1     i   .req r2     t   .req r3     ret .req r0      mov i, #0      ; Unoptimized loop needs 11 cycles loop:     ldrb t, [src, +i]     cmp t, #0     beq loop_end     cmp t, c     beq found     add i, i, #1     b loop found:     mov ret, i     b end loop_end:     mov ret, #-1     b end end:     bx lr</pre>	<pre>.global my_strchr  my_strchr:     src .req r0     c   .req r1     i   .req r2     t   .req r3     ret .req r0      mov i, #0      ; t is preloaded     ldrb t, [src, +i]      ; Optimized loop needs 9 cycles loop:     cmp t, #0     beq loop_end     cmp t, c     beq found     add i, i, #1      ; t is preloaded     ldrb t, [src, +i]     b loop      ; Removed two useless branches below found:     mov ret, i     bx lr loop_end:     mov ret, #-1     bx lr</pre>

The existing code has a number of inefficiencies that need to be fixed. First, there is a two cycle stall after the `ldrb` instruction because of the dependency on `t`. Second, there are unnecessary `b end` instructions involved in the end of the routine that can be removed.

To fix the stalls, pre-loading of `t` is used. This means that there will be a 2 cycle stall for the first iteration of the loop, but future iterations will not have a stall. For most loop runs, this will result in a 2 cycle savings. The original loop requires 11 cycles in the normal case, and the optimized code requires only 9 cycles. To fix the unnecessary `b end` instructions, the `end` label is removed and `bx lr` is repeated (called inlining) in both places where `b end` was previously. This will save 3 cycles as the subroutine finishes, and also removes one instruction from the code.

I decided to apply these changes because they are most straightforward approach to optimizing this code. I also investigated the possibility of applying partial loop unrolling (for example, checking 4 characters per loop iteration) but if you do that then you need to add extra code to the loop to properly calculate `i`, which removes the efficiency gain.

## strcmp

Original	Optimized
<pre> .global my_strcmp  my_strcmp:     str1 .req r0     str2 .req r1     i .req r2     t1 .req r3     t2 .req r4     ret .req r0      stmfd sp!,{r4}      mov i, #0     ; Unoptimized loop needs 12 cycles loop:     ldrb t1, [str1, +i]     ldrb t2, [str2, +i]     cmp t1, t2     bne fail     cmp t1, #0     beq loop_end     add i, i, #1     b loop fail:     mov ret, i     b end loop_end:     mov ret, #0     b end end:     ldmfd sp!, {r4}     bx lr </pre>	<pre> .global my_strcmp  my_strcmp:     str1 .req r0     str2 .req r1     i .req r2     t1 .req r3     t2 .req r4     ret .req r0      stmfd sp!,{r4}      mov i, #0     ; t1 and t2 are preloaded     ldrb t1, [str1, +i]     ldrb t2, [str2, +i]     ; Optimized loop needs 10 cycles loop:     cmp t1, t2     bne fail     cmp t1, #0     beq loop_end     add i, i, #1      ; t1 and t2 are preloaded     ldrb t1, [str1, +i]     ldrb t2, [str2, +i]      b loop     ; Removed two useless branches below fail:     mov ret, i     ldmfd sp!, {r4}     bx lr loop_end:     mov ret, #0     ldmfd sp!, {r4}     bx lr </pre>

The existing code has a number of inefficiencies that need to be fixed. First, there is a one cycle stall after the `ldrb t1` instruction and another two cycle stall from the `ldrb t2` instruction. (However, the one cycle stall becomes part of the two cycle stall because of parallelism in the pipeline, so the net stall is two cycles.) Second, there are multiple unnecessary branches at the end of the program.

To fix the stalls, pre-loading of both `t1` and `t2` is used. This means that there will be a 2 cycle stall for the first iteration of the loop, but future iterations will not have a stall. For most loop runs, this will result in a 2 cycle savings. To fix the unnecessary branches, the branch structure is removed and conditional execution is added. This will save 3 cycles as the subroutine finishes, as only one of those branches is taken at a time.

I decided to apply these changes because they are most straightforward approach to optimizing this code. I also investigated the possibility of applying partial loop unrolling, but if you do that then you will quickly need lots of extra registers.

## **Conclusion**

In this report, I presented the optimization of five different string library functions written in ARM assembly. The primary optimizations employed were preloading to remove stalls and removing branches by using conditional instructions. I could have done more optimization to some of the functions by applying loop unrolling, but in some cases it wasn't feasible.