

SPRING 2016

CMPE 364

Microprocessor Based Design

Dr. Mohamed Al-Meer

Lecture Expectations

- Expected to achieve:
 - Writing assembly code
 - Cycle Counting
 - Instruction scheduling
 - Load instructions
 - Preloading
 - Loop unrolling
 - Register Allocation
 - Allocating variables to registers
 - Making the most of available registers

Writing Assembly Code

- Embedded software projects contain a few key subroutines that dominate system performance.
 - By **optimizing** these routines you can **reduce the system power consumption**
 - **Reduce the clock speed** needed for real-time operation.
- Optimization can turn an
 - infeasible system into a feasible one, or
 - an uncompetitive system into a competitive one.

Writing Assembly Code

- **For maximum performance**, you can optimize critical routines using **hand-written assembly**.
- **SO** we will address in this chapters three major Optimization Tools
 - ***Instruction scheduling***: Reordering the instructions in a code sequence to avoid processor stalls.
 - ***Register allocation***: Deciding how variables should be allocated to ARM registers or stack locations for maximum performance
 - ***Conditional execution***: Accessing the full range of ARM condition codes and conditional instructions

Instruction Scheduling

- The time taken to execute instructions depends on the implementation pipeline.
 - (Assume **ARM9TDMI** pipeline timings)
- Instructions that are conditional on the value of the ARM condition codes in the *cpsr* take **one cycle** if the **condition is not met**.
- But if **met** then the following rules apply:

Instruction Scheduling

- **ALU** operations such as addition, subtraction, and logical operations take **one cycle**.
- **LOAD** instructions that load ***N* 32-bit** words of memory such as LDR and LDM take ***N* cycles** to issue
 - but the result of the last word loaded **is not available on the following one cycle**. updated load address is available on the next cycle
- **LOAD** instructions that load **16-bit** or **8-bit** data such as LDRB, LDRSB, LDRH, and LDRSH take **one cycle** to issue.
 - The load result is **not available** on the following **two cycles**.

Instruction Scheduling

- **Branch** instructions take **three** cycles.
- **Store** instructions that store ***N* values** take ***N* cycles**.
 - An STM of a single value is exceptional, taking two cycles.
- **Multiply** instructions take a **varying** number of **cycles** depending on the value of the second operand in the product (see Table D.6 in Section D.3).

ARM9TDMI Processor Pipeline

- **Fetch**: Fetch from memory the instruction at address *pc*.
- **Decode**: Decode the instruction that was fetched in the previous cycle.
 - The processor also **reads the input operands from the register bank** if they are **not available via one of the forwarding paths**.
- **ALU**: Executes the instruction that was decoded in the previous cycle.
 - Note this instruction was originally fetched from address ***pc* – 8 (ARM state)** or ***pc* – 4 (Thumb state)**.
 - Some instructions may spend several cycles in this stage.
 - Or do **addressing calculations**.

ARM9TDMI Processor Pipeline

- **LS1**: Load or store the data specified by a load or store instruction.
 - If the instruction is not a load or store, then this stage has no effect.
- **LS2**: Extract and zero- or sign-extend the data loaded by a byte or half word load instruction.
 - If the instruction is not a load of an 8-bit byte or 16-bit half word item, then this stage has no effect.

Instruction address	<i>pc</i>	<i>pc-4</i>	<i>pc-8</i>	<i>pc-12</i>	<i>pc-16</i>
Action	Fetch	Decode	ALU	LS1	LS2

Example - 1

- This example shows the case where there is no interlock.
 - ADD r0, r0, r1**
 - ADD r0, r0, r2**
- This instruction pair takes **two cycles**. The ALU calculates $r0 + r1$ in one cycle
 - result is available for the ALU to calculate $r0 + r2$ in the second cycle.
- **Because it uses Forwarding found in ARM9TDMI.**

Example - 2

Pipeline	Fetch	Decode	ALU	LS1	LS2
Cycle 1	...	ADD	LDR	...	
Cycle 2		...	ADD	LDR	...
Cycle 3		...	ADD	—	LDR

- This example shows a **one-cycle interlock** caused by load use
 - LDR *r1*, [r2, #4]** 2 cycles
 - ADD r0, r0, *r1*** 1 cycle
- Takes **three** cycles.
- ALU calculates the address $r2 + 4$ in the first cycle while decoding the ADD instruction in parallel.
- Pipeline stalls for one cycle.
- Processor executes the ADD in the ALU on the third cycle

Example - 3

- This example shows a one-cycle interlock caused by delayed load use.
 - LDRB *r1*, [r2, #1]** 3 Cycles
 - ADD r0, r0, r2** 1 Cycle (within LDRB!)
 - EOR r0, r0, *r1*** 1 Cycle
- Takes **four** cycles.
- When EOR starts at cycle 3, LDRB need another cycle to write back r1.
- So ARM stalls EOR for 1 cycle more.

Scheduling of load instructions

- Load instructions occur frequently in compiled code, accounting for one third of all instructions!!
 - **Very important? Why?**
 - **Make it faster >>makes execution faster by 30%.**
- Careful scheduling of load instructions so that **pipeline stalls don't occur can** improve performance.
- The compiler attempts to schedule the code as best it can.
 - But still **needs user optimization.**

Scheduling of load instructions Example

- Let's consider an example of a memory-intensive task.
- The following function, *str_tolower*, copies a zero-terminated string of characters from *in* to *out*. It converts the string to lowercase in the process.
- **See next example in C and ARM assembly**

Scheduling of load instructions Example

```
void str_tolower(char *out, char *in)
{
    unsigned int c;

    do
    {
        c = *(in++);
        if (c>='A' && c<='Z')
        {
            c = c + ('a' - 'A');
        }
        *(out++) = (char)c;
    } while (c);
}
```

Compiler Code

```
str_tolower
    LDRB    r2,[r1],#1    ; c = *(in++)
    SUB     r3,r2,#0x41    ; r3 = c - 'A'
    CMP     r3,#0x19      ; if (c <='Z'-'A')
    ADDLS   r2,r2,#0x20    ; c += 'a'-'A'
    STRB    r2,[r0],#1    ; *(out++) = (char)c
    CMP     r2,#0          ; if (c!=0)
    BNE     str_tolower    ; goto str_tolower
    MOV     pc,r14         ; return
```

Scheduling of load instructions Example

- Unfortunately, the **SUB** instruction uses the value of **c** directly after the **LDRB** instruction that loads c. Consequently, the ARM9TDMI pipeline will **stall for two cycles**.
- Everything following the load of c depends on its value!
- But
 - **two ways** you can alter the structure of the algorithm to avoid the cycles by using assembly
- This takes **11 cycles** to execute. See how 11??

Load Scheduling by Preloading

```

out    RN 0    ; pointer to output string
in     RN 1    ; pointer to input string

c      RN 2    ; character loaded
t      RN 3    ; scratch register
; void str_tolower_preload(char *out, char *in)
str_tolower_preload
LDRB   c, [in], #1      ; c = *(in++)

loop
    SUB    t, c, #'A'      ; t = c-'A'
    CMP    t, #'Z'-'A'     ; if (t <= 'Z'-'A')
    ADDLS  c, c, #'a'-'A'   ; c += 'a'-'A';
    STRB   c, [out], #1    ; *(out++) = (char)c;
    TEQ    c, #0           ; test if c==0
    LDRNEB c, [in], #1     ; if (c!=0) { c=*(in++);
    BNE    loop           ;             goto loop; }
    MOV    pc, lr         ; return

```

- we load the data required for the loop at the end of the previous loop.
- To get performance improvement with little increase in code size.
- one instruction longer than the C version, but we save two cycles for each inner loop iteration.
- reduces the loop from **11 cycles** per character to **9 cycles** per character on an ARM9TDMI, giving a **1.22 times** speed improvement

Load Scheduling by Unrolling

- Works by unrolling and then interleaving the body of the loop.
 - For Example: loop iterations i , $i + 1$, $i + 2$ interleaved.
- When the result of an operation from loop i is not ready, we can perform an operation from loop $i + 1$.

Load Scheduling by Unrolling Example

- The next program applies **load scheduling** by **unrolling** to the *str_tolower* function.
- **See**
 - Very efficient
 - Requires **7 cycles** per character on ARM9TDMI.
 - **1.57 times** speed increase over the original *str_tolower*.
 - We use **conditional instructions** to avoid storing characters that are past the end of the string.
 - Improvements cost: The routine is more than **double** the code size of the original implementation
 - For **time-critical** parts of an application where you know the data size is large

```

out    RN 0    ; pointer to output string
in     RN 1    ; pointer to input string
ca0    RN 2    ; character 0
t      RN 3    ; scratch register

ca1    RN 12   ; character 1
ca2    RN 14   ; character 2
; void str_tolower_unrolled(char *out, char *in)
str_tolower_unrolled
STMFD  sp!, {lr}      ; function entry
loop_next3
LDRB   ca0, [in], #1    ; ca0 = *in++;
LDRB   ca1, [in], #1    ; ca1 = *in++;
LDRB   ca2, [in], #1    ; ca2 = *in++;
SUB    t, ca0, #'A'      ; convert ca0 to lower case
CMP    t, #'Z'-'A'
ADDLS  ca0, ca0, #'a'-'A'
SUB    t, ca1, #'A'      ; convert ca1 to lower case
CMP    t, #'Z'-'A'
ADDLS  ca1, ca1, #'a'-'A'
STRB   ca0, [out], #1    ; *out++ = ca0;
TEQ    ca0, #0           ; if (ca0!=0)
STRNEB ca1, [out], #1    ; *out++ = ca1;
TEQNE  ca1, #0           ; if (ca0!=0 && ca1!=0)
STRNEB ca2, [out], #1    ; *out++ = ca2;
TEQNE  ca2, #0           ; if (ca0!=0 && ca1!=0 && ca2!=0)
BNE    loop_next3       ; goto loop_next3;
LDMFD  sp!, {pc}        ; return;

```

Register Allocation

- You can use **14** of the 16 visible ARM registers to hold **general-purpose data**.
- For a function to be AAPCS compliant it must preserve the callee values of registers **r4 to r11**.
- Use the following template for optimized assembly routines requiring many registers:

Register Allocation

```
routine_name
    STMFD sp!,      {r4-r12, lr}      ; stack saved registers
    ; body of routine
    ; the fourteen registers r0-r12 and lr are available
    LDMFD sp!,      {r4-r12, pc}      ; restore registers and return
```

```
routine_name
    STMFD sp!,      {r4-r12, lr}      ; stack saved registers
    ; body of routine
    ; registers r0-r12 and lr available
    LDMFD sp!,      {r4-r12, lr}      ; restore registers
    BX              lr                ; return, with mode switch
```